

Per-model accuracy: the codec is consistent; LLMingua-2 swings 0.25--1.00

