

Framework scorecard: the codec is the only method green on every correctness axis

method	tokens saved	legible?	lossless?	facts kept	end-to-end accuracy	acc spread (5 models)
recency	64%	yes	no	12%	0.00	0.00-0.00
digest (AFM)	67%	yes	no	0%	0.00	0.00-0.00
LLMLingua-2 (matched)	80%	yes	no	100%	0.80	0.25-1.00
LLMLingua-2 (aggressive)	94%	yes	no	62%	--	--
codec (ours)	75%	yes	yes	100%	0.95	0.75-1.00

In-context methods (tokens). Byte codecs (gzip/zstd/brotli) are a separate storage class: higher raw ratio but not legible (Fig. compare b).